## Notes on Principle Component Analysis

## Nathaniel

I have long wanted to dive into the intricacy of Principle Component Analysis (PCA). This note serves as my attempt to do so. In the following, I will explore what PCA is and apply it on interest rate dynamics as an example using Matlab, Stata, R and SAS.

PCA is one of the most popular techniques in unsupervised learning where only data on independent variables (features) are given. The key idea behind PCA, is to use orthogonal projections to find lower dimensional representations of data that retain as much information as possible. It is also often used as a tool in exploratory data analysis and as a dimension reduction technique.

To motivate the study of PCA, I set out to apply PCA on interest rate dynamics. The dataset, taken from Miller (2013), is daily U.S. government rates from March 2000 through August 2000 with six variables representing maturities of 1, 2, 3, 5, 10, and 30 years. I hope to demonstrate the following observation made by Dowd (2005) P.120

"The standard financial example is where the original variables might be different spot (or interest) rates across the maturity spectrum, and where the first three principal components are reported to explain over 95% of spot-rate behaviour... the first principal components can be interpreted as reflecting the level of the spot-rate curve, the second can be interpreted as reflecting its steepness; and the third can be interpreted as reflecting its curvature. We can therefore model the spot rates using only a small number of spot-rate principal components—which shows that PCA can be useful for reducing the dimensionality of a multivariate problem."

Figure 1 recreates the example from Miller (2013) Exhibits 9.17 (P.190) and Excel workbook (Chapter\_09\_PCA\_Interest\_Rates). It shows that using only the first three of the six principle component, one can approximate the actual on-year yield curve very closely. PCA can thus serves "as a basis for an interest rate model or as the basis for a risk report. A portfolio's correlation with these PCs might also be a meaningful risk metric." (Miller 2013, Page 189).

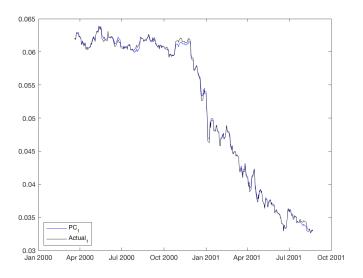


Figure 1: Actual and Approximate 1 Rates (Miller, 2013, Exhibit 9.17)

## The Mathematical Idea behind PCA

Before diving into PCA, it would be best to review the following concepts in linear algebra:

- Inner products for computing the lengths, distances, and angles between vectors. An inner product is defined as a symmetric, positive definite, bilinear mapping. An example would be  $\beta(x^T A y)$ .
- The big picture (from MIT 1806) of linear algebra—column space, left null space, row space and null space from the equation Ax = b.
- Orthogonal projections of data onto lower dimensional subspaces. Recall  $P_X = X(X^TX)^{-1}X^T$  and  $M_X = I P_X = I X(X^TX)^{-1}X^T$ .
- Orthonormal basis. The orthogonal square matrix Q and it properties such as  $Q^TQ = I$ ,  $Q^T = Q^{-1}$ . Gram-Schmidt process. QR decomposition to decompose any real square matrix A into Q and an upper triangular matrix R. The projection matrix becomes  $QQ^T$ . The projection solution is found from  $Q^TQx = Q^Tb$ .
- Eigendecomposition (Spectral Decomposition). Covariance matrix as positive definite symmetric matrix (all eigenvalues are positive and all pivots are positive).  $A = UDU^T$  (Rotate Strech Rotate), where D is a diagonal matrix. The entries of the diagonal of D are the eigenvalues of A. The column vectors of U are the eigenvectors of A and they are orthonormal.
- Singular value decomposition (SVD). Spectral Decomposition is one of the special cases of SVD where U = V = Q if A is symmetric positive definite. SVD transform a set of orthogonal bases in one space (row space) to another (column space) while keeping orthogonality of the two sets of bases intact. The SVD theorem states:

$$A_{nxp} = U_{n \times n} S_{n \times p} V_{n \times n}^T$$

where  $U^TU = I_{n \times n}$  and  $V^TV = I_{p \times p}$  (i.e. U and V are orthogonal; the columns of U are the left singular vectors; S has singular values and is diagonal; and  $V^T$  has rows that are the right singular vectors. The SVD represents an expansion of the original data in a coordinate system where the covariance matrix is diagonal.

Calculating the SVD consists of finding the eigenvalues and eigenvectors of  $AA^T$  and  $A^TA$ . The eigenvectors of  $A^TA$  make up the columns of V, the eigenvectors of  $AA^T$  make up the columns of U. Also, the singular values in S are square roots of eigenvalues from  $AA^T$  or  $A^TA$ . The singular values are the diagonal entries of the S matrix and are arranged in descending order. The singular values are always real numbers. If the matrix A is a real matrix, then U and v are also real.

PCA is simply an application of the SVD. The principal components are equal to the right singular values if you first scale the data by subtracting the column mean and dividing each column by its standard deviation (that can be done with the scale() function). Using the SVD, the score matrix T is:  $T = AX = USV^TV = US$  and one can use the first k column of T (where k < m) for dimensionality reduction purpose.

There are at least five different perspectives of PCA that lead to different objectives:

- 1. minimising the squared reconstruction error,
- 2. minimising the auto-encoder loss,
- 3. maximising the mutual information,
- 4. maximising the variance of the projected data, and
- 5. maximising the likelihood in a latent variable model

All these different perspectives give us the same solution to the PCA problem. The strengths and weaknesses of individual perspectives become more clear and important when we consider properties of real data. This note focuses on the first and third perspectives.

So what is PCA? In a given vector space, there is potentially an infinite number of orthonormal bases. The question is whether there are any orthonormal bases more 'special' than others? In PCA, a basis is chosen so that the first vector in the basis (principal component), explains as much of the variance in the data, as possible. In higher dimensions, each successive principal component explains the maximum amount of variance in the residual data after taking into account all of the preceding components (orthogonal complement to its 'previous' principle subspace). Just as the first principal component explained as much of the variance in the data as possible, the second principal component explains as much of the variance in the residuals and so on while keeping the covariance between principle components zero.

Borrowing from the notion in Tsay (2010) P.483 and Dowd (2005) P.119, the above paragraph on maximising the variance of the projected data can be translated to the following:

$$max_w : Var(y_i) = \mathbf{w_i^T} \mathbf{\Sigma_r} \mathbf{w_i}$$
  $i = 1, ..., k,$   
 $min_w : Cov(y_i, y_j) = \mathbf{w_i^T} \mathbf{\Sigma_r} \mathbf{w_i} = 0,$   $i \neq j$ 

subject to the constraints that the length of the orthonormal basis is 1, that is  $w_i^T w_i = \sum_{j=1}^k w_{ij}^2 = 1$ . The problem at hand could be solved by Spectral Decomposition. and the solution is that:

The *i*th principal component of **r** is  $y_i = e_i' r = \sum_{j=1}^k e_{ij} r_j$  for i = 1, ..., k. Moreover,

$$Var(y_i) = e_i' \Sigma_r e_i = \lambda_i,$$
  $i = 1, ..., k,$   
 $Cov(y_i, y_j) = e_i' \Sigma_r e_i = 0,$   $i \neq j$ 

To understand how the above equations are set up, let us dive into some more mathematics. Given data centred with mean zero and represented by matrix  $X_n$  with D column of  $x_i...x_d$ ,  $X_n$  can be represented by orthonormal basis matrix B which consists of orthonormal bases  $b_i...b_D$ .

$$X_n = \sum_{i=1}^D \beta_{in} b_i$$

$$X_n = \sum_{i=1}^{M} \beta_{in} b_i + \sum_{i=M+1}^{D} \beta_{in} b_i$$

The key idea in PCA is to find a lower dimensional representation of  $X_n$ , denote as  $\tilde{X}_n$  that can be expressed using fewer basis vectors.  $\tilde{X}_n$  is still a D-dimensional vector, it lives in an M-dimensional subspace of  $\mathbb{R}^{\mathbb{D}}$  and only M coordinates:  $\beta_{n1}...\beta_{n_M}$  are necessary to represent it, which leads to the deletion of the second component

$$\tilde{X}_n = \sum_{i=1}^M \beta_{in} b_i$$

We want to find parameters  $\beta_{in}$  and orthonormal basis vectors  $b_i...b_D$  which span the principle subspace, such at the loss function—the average squared reconstruction error is minimised.  $\beta_{in}$  can be interpreted as weights of the orthonormal bases. But if one was allowed to make those weights as big as he/she wanted, one could make the variance of as big as possible. The natural constraint is to normalize  $\beta_{in}$  so the sum of their squares is 1. The average squared reconstruction error is denoted as J optimized through standard derivative procedure (including multivariate chain rule). The whole procedure is detailed in the online course, Mathematics for Machine Learning: PCA, offered by the Imperial College London (week 4). The key equations to note are:

$$min_{\tilde{x}_n}: J = \frac{1}{N} \sum_{n=1}^{N} ||x_n - \tilde{x}_n||^2$$

$$\frac{\partial J}{\partial \beta_i n} = \frac{\partial J}{\partial \tilde{x}_n} * \frac{\partial \tilde{x}_n}{\partial \beta_i n}$$
$$\frac{\partial J}{\partial \beta_i n} = -\frac{2}{N} (x_n - \tilde{x}_n)^T * b_i = 0$$

$$\Rightarrow \beta_{in}^* = x_n^T b_i$$

Substituting  $\beta_{in}^* = x_n^T b_i$  into  $x_n - \tilde{x}_n$  and reformulating the J gives

$$J = \sum_{j=M+1}^{D} b_{J}^{T} S b_{j} = trace \left( \sum_{j=M+1}^{D} b_{i} b_{j}^{T} \right) S$$

where S is the data covariance matrix.

In a 2 dimensional example, the Lagrangian would be set up as follow:

$$L = b_2^T S b_2 + \lambda (1 - b_2^T b_2)$$

Setting the partial derivatives equal to zero and solving for solutions yield:

$$J = \lambda$$

Using Lagrange multipliers generally, it could be shown that

$$J = \sum_{j=M+1}^{D} \lambda_j$$

and

$$\tilde{X} = B \underbrace{B^T x}_{coordinate}$$

In other words, J, the average reconstruction error is minimised if we choose the <u>basis vectors</u> that span the ignored subspace to be the <u>eigenvectors</u> of the data covariance matrix that belong to the smallest eigenvalues. This equivalently means that the principal subspace is spanned by the eigenvectors belonging to the M largest eigenvalues of the data covariance matrix. If the basis vectors are picked in this way, the reconstruction error is minimized and equal to the sum of D - M - 1 smallest eigenvalues of the (full) data covariance matrix.

This nicely aligns with properties of the covariance matrix. The eigenvectors of the covariance matrix are orthogonal to each other because of symmetry and the eigenvector belonging to the largest eigenvalue points in the direction of the data with the largest variance and the variance in that direction is given by the corresponding eigenvalue. Similarly, the eigenvector belonging to the second largest eigenvalue points in the direction of the second largest variance of the data and so on.

In summary, PCA involves the following steps:

- 1. Standardize the raw data (correlation is still intact).
- 2. Calculate a covariance matrix of the standardized data.
- 3. Decompose the covariance matrix to obtain eigenvalues and eigenvectors.

4. Project data onto the principal subspace spanned by the first n eigenvectors ranked according to their eigenvalues (largest to smallest).

A simple 2-vector example in Matlab shows that PC linearly transform data:

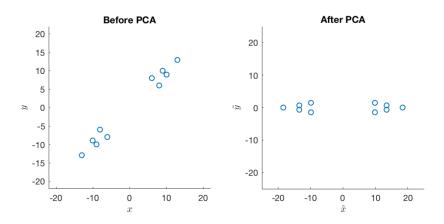


Figure 2

Figure 1 (right) shows that the largest PC (x) spans the direction that maximizes the variance of the projected data, and the smallest PC (y) maximizes that variance of the residual. In effect, PCA linearly transforms the data points we had. Quoting from ISLR P.379

The appeal of this interpretation is clear: we seek a single dimension of the data that lies as close as possible to all of the data points, since such a line will likely provide a good summary of the data ... The notion of principal components as the dimensions that are closest to the n observations extends beyond just the first principal component. For instance, the first two principal components of a data set span the plane that is closest to the n observations, in terms of average squared Euclidean distance.

## An Example using PCA on Interest Rate Dynamic

Let us return to the interest rate dynamic example from Miller (2013).

From Matlab (see textttPCA\_Practice.m), we can derive the eigenvalues as below

$$\lambda_1$$
  $\lambda_2$   $\lambda_3$   $\lambda_4$   $\lambda_5$   $\lambda_6$  5.39 0.56 0.041 0.004 0.002 0.0006

the scoring coefficients (or equivalently loading, orthonormal eigenvectors) matrix is:

0.3910	-0.5335	0.6101	0.3367	0.2260	-0.1602
0.4220	-0.2630	-0.03014	-0.3087	-0.2675	0.76476
0.4268	-0.1631	-0.1981	-0.3562	-0.4949	-0.6164
0.42853	0.01134	-0.4604	-0.1798	0.7538	-0.0595
0.4186	0.2949	-0.3152	0.7555	-0.2486	0.0760
0.3576	0.7296	0.5255	-0.2473	0.0469	-0.00916

with the constraint that the norms of eignvector is equal to 1 (for instance, the sum of squares of the first eigenvector is:  $0.391^2 + 0.4221^2 + 0.4269^2 + 0.4286^2 + 0.4187^2 + 0.3574^2 \approx 1.00003543$ )

And the proportions of the total variance explained by each principle component are

$$\lambda_1/\sum \lambda$$
  $\lambda_2/\sum \lambda$   $\lambda_3/\sum \lambda$   $\lambda_4/\sum \lambda$   $\lambda_5/\sum \lambda$   $\lambda_6/\sum \lambda$   $0.89837$   $0.99198$   $0.99884$   $0.99953$   $0.99989$   $1.0000$ 

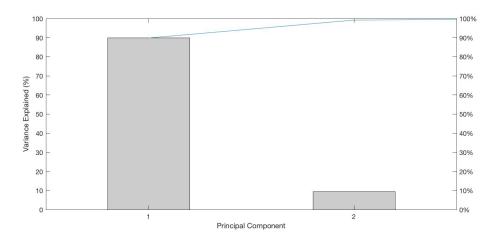


Figure 3: The first two PCs explains over 90% of the variance.

In essence, this example demonstrates the prototypical patterns-shifting, tilting, and twisting seen in PCA. As note by Miller (2013) P189-190:

[The first PC] This flat, equal weighting represents the shift of the yield curve. A movement in this component increases or decreases all of the points on the yield curve by the same amount (actually, because we standardized all of the data, it shifts them in proportion to their standard deviation). Similarly, the second PC shows an upward trend. A movement in this component tends to tilt the yield curve. Finally, if we plot the third PC, it is bowed, high in the center and low on the ends. A shift in this component tends to twist the yield curve ...

We could multiply all of the elements in one column of the PC matrix ... The justification for doing this is purely aesthetic ...

Because the first three PCs explain so much of the dynamics of the yield curve, they could serve as a basis for an interest rate model or as the basis for a risk report. A portfolio's correlation with these PCs might also be a meaningful risk metric.

which echos Dowd (2005)'s "level of the spot-rate curve, steepness ( $\sim$  duration); curvature ( $\sim$  convexity)".

Note that PCA is closely related to Singular value decomposition (SVD). The principal components are equal to the right singular values if you first scale (subtract the mean, divide by the standard deviation) the variables. For instance, Dowd (2005)'s code (such as PCAVAR and PCAELT) use the SVD instead. The code above calculates the differences between the two methods and confirms the result (see comment).

The loading plot examines the first two PC loadings for the daily U.S. government rates dataset. To see the relationship between loading and eignevector, see https://stats.stackexchange.com/questions/143905/loadings-vs-eigenvectors-in-pca-when-to-use-one-or-another. From the first component, we see that U.S. government rate with 30-year maturity is separated from others. Next comes the 1-year maturity which is in the "middle", potentially hightlighting the differences of the two variables from others. From the second component, the difference between the government rate with 1-year and the 30 year of maturity again stands out. Recall that the 2nd PC could be interpreted as steepness of the yield curve. This can be seen from the ascending order of the variables from -0.5 to 0.8 (Componmet 2). Also note that ordering of PC2 is very similar to that of eignvector2 in Figure 2. This follows the defintion of loadings.

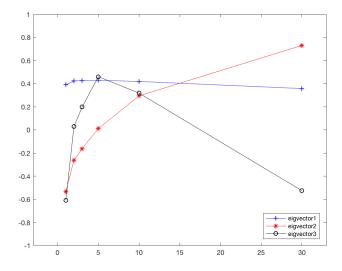


Figure 4: First Three Principal Components' loadings of the Yield Curve (Miller, 2013, Exhibit 9.16)

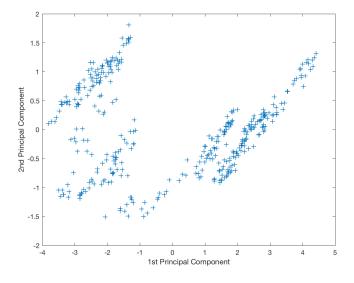


Figure 5: Actual and Approximate 1 Rates (Miller, 2013, Exhibit 9.17)

Reference Dowd Kevin (2005) Measuring Market Risk, Wiley, 2nd Edition

Ruey S. Tsay (2010) Analysis of Financial Time Series, Wiley, 3th Edition

G.James, D.Witten, T.Hastie and R.Tibshirani (2013), An Introduction to Statistical Learning, with applications in R (ISLR), Springer

Michael B. Miller (2013) Mathematics and Statistics for Financial Risk Management, Wiley, 2nd Edition, ISBN: 978-1-118-75029-2

Hull, John. (1989). Options, futures, and other derivative securities. Englewood Cliffs, N.J.: Prentice Hall.

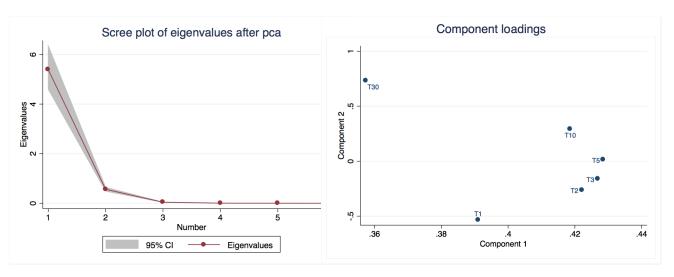


Figure 6 Figure 7